



Lurkr

always watching, never blinking

Multi-agent market intelligence — the competitor-watching team that never sleeps.

Founders fly blind on competition

- ▶ Competitor research is manual, scattered across tabs — and stale the day it's done.
- ▶ Generic “market reports” aren't about your product or your specific edge.
- ▶ By the time you notice a rival's move, the window to react has already closed.

Describe your idea. Get a personalised brief.

- ▶ You describe your startup in a sentence; Lurkr finds the real competitors.
- ▶ A team of AI agents gathers live signals and analyses them for YOUR product.
- ▶ Output: the single biggest threat + opportunity, each with a recommended action.
- ▶ And it keeps watching — automatic daily refresh while you sleep.

A grounded multi-agent pipeline

- ▶ ① Discovery — finds real, currently-operating competitors in your space.
- ▶ ② Gather — live web (Tavily) + news (Google News) signals per competitor.
- ▶ ③ Analysts — Marketing · Product · Sales run in parallel, grounded in those signals.
- ▶ ④ Strategy — synthesises one executive brief: threat + opportunity + watch items.

Tiered, secret-isolating topology

- ▶ Thin-client Android shell (Capacitor) loads a CDN-hosted React 19 SPA.
- ▶ Stateless Express API tier holds every key — nothing secret ships on the client.
- ▶ Neon Postgres stores accounts, ideas and cached analyses.
- ▶ OpenRouter (LLMs) · Tavily + Google News (signals) · GitHub Actions (cron + CI/CD).

Why the multi-agent design wins

- ▶ Fan-out / fan-in: three analysts run concurrently, Strategy synthesises the result.
- ▶ Grounding / RAG discipline: every claim must cite a real signal — anti-hallucination.
- ▶ Cost-aware model tiering: a fast model for breadth, a stronger one for judgement.
- ▶ Fault-isolated sources + bounded-retry JSON parsing — the sweep always completes.

Why this stack

- ▶ Neon (serverless Postgres) — relational + JSONB, scale-to-zero, DB branching, zero lock-in.
- ▶ OpenRouter — one OpenAI-compatible API across models; cost-aware tiering, no lock-in.
- ▶ Roll-our-own auth — script + HMAC JWT on the Node std-lib; no auth SDK, data stays ours.
- ▶ Render + GitHub Pages + Actions — free-tier, zero-ops, ship by git push.

Per-user intelligence that compounds

- ▶ Hard auth gate — Google (native account picker on Android) or email + password.
- ▶ Every search is saved per user; a “My Ideas” library keeps them.
- ▶ Stale-while-present cache: instant re-opens; a live sweep runs only on a cache miss.
- ▶ Opt-in daily refresh at 04:00 IST — a cron re-runs and caches your watched ideas.

Build-once, ship by git push

- ▶ The Android APK is a thin shell — frontend updates need no rebuild or reinstall.
- ▶ A signed, auto-versioned APK pipeline publishes to GitHub Releases.
- ▶ Frontend auto-deploys to GitHub Pages; the backend auto-deploys on Render.
- ▶ A CI build-gate compiles the app on every change before it ships.

Security & resilience by default

- ▶ Secret isolation — keys live only on the backend; the client bundle ships none.
- ▶ Deny-by-default gate; the user is revalidated from the DB on every request.
- ▶ Memory-hard scrypt hashing, constant-time verify, audience-pinned Google tokens.
- ▶ Idempotent migrations; an append-only usage ledger ready for usage-based billing.

What's next

- ▶ Usage-based billing — the usage ledger is already in place.
- ▶ Email alerts when a new high-urgency threat appears (Resend).
- ▶ Email verification + password reset, once a sending domain is provisioned.
- ▶ Native Google sign-in polish and an iOS build.



Lurkr

Always watching, never blinking.

React 19 · Vite · Tailwind v4 · Capacitor 7 · Express · Neon Postgres · OpenRouter · GitHub Actions · Render